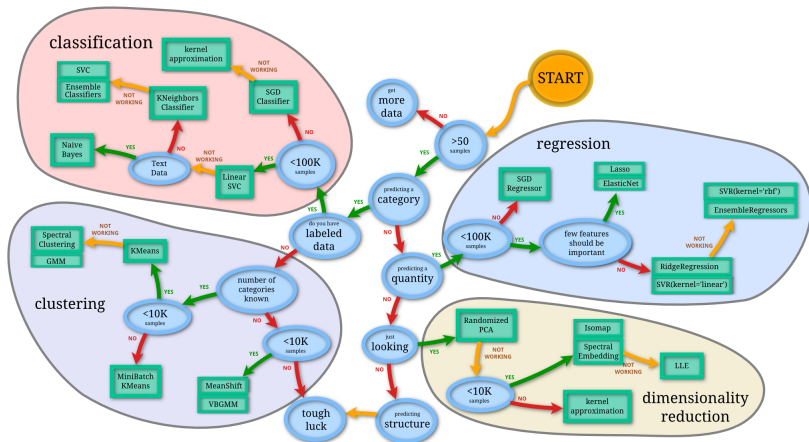


PATTERN RECOGNITION

Bertrand Thirion and John Ashburner



SOME KEY CONCEPTS

supervised learning: the data comes with additional attributes that we want to predict \implies classification and regression.

unsupervised learning: no target values.

- discover groups of similar examples within the data (clustering)
- determine the distribution of data within the input space (density estimation)
- project the data down to two or three dimensions for visualization

GENERAL SETTING

We have a training dataset of n observations, each consisting of an input \mathbf{x}_i and a target y_i .

Each input, \mathbf{x}_i , consists of a vector of p features.

$$\mathcal{D} = \{(\mathbf{x}_i, y_i) | i = 1, \dots, n\}$$

The aim is to predict the target for a new input \mathbf{x}_* .


CLASSIFICATION

Targets (\mathbf{y}) are categorical labels.

Train with

$$\mathcal{D} = \{(\mathbf{x}_i, y_i) | i = 1, \dots, n\}$$

and use result to make best guess of y_* given \mathbf{x}_* .



PROBABILISTIC CLASSIFICATION

Targets (\mathbf{y}) are categorical labels.


Train with

$$\mathcal{D} = \{(\mathbf{x}_i, y_i) | i = 1, \dots, n\}$$

and compute

$$P(y_* = 1 | \mathbf{x}_*, \mathcal{D}).$$

Probabilistic classification




REGRESSION

Targets (\mathbf{y}) are continuous real variables.

Train with

$\mathcal{D} = \{(\mathbf{x}_i, y_i) | i = 1, \dots, n\}$
and compute $p(y_* | \mathbf{x}_*, \mathcal{D})$.



MANY OTHER SETTINGS

- **Multi-class classification** when there are more than two possible categories.
- **Ordinal regression** for classification when there is some ordering of the categories.
Chu, Wei, and Zoubin Ghahramani. "Gaussian processes for ordinal regression." In Journal of Machine Learning Research, pp. 1019-1041. 2005.
- **Multi-task learning** when there are multiple targets to predict, which may be related.
- etc


MULTI-CLASS CLASSIFICATION

- **Multinomial Logistic regression** Theoretically optimal.
Expensive optimization. rarely used in practice.
- **One-versus-all classification [SVMs]** Among several hyperplane, choose the one with maximal margin. → recommended
- **One-versus-one classification** Vote across each pair of class.
Expensive, not optimal.

CURSE OF DIMENSIONALITY



Large p , small n .

NEAREST-NEIGHBOUR CLASSIFICATION






- Not nice smooth separations.
- Lots of sharp corners.
- May be improved with *K-nearest neighbours*.

BEHAVIOUR CHANGES IN HIGH-DIMENSIONS



BEHAVIOUR CHANGES IN HIGH-DIMENSIONS



OCCAM'S RAZOR

"Everything should be kept as simple as possible, but no simpler."

— Einstein (allegedly)

- Complex models (with many estimated parameters) usually explain training data better than simpler models.
- Simpler models often generalise better to new data than more complex models.

Need to find the model with the optimal bias/variance tradeoff.

BAYESIAN MODEL SELECTION

Real Bayesians don't cross-validate (except when they need to).

$$P(M|\mathcal{D}) = \frac{p(\mathcal{D}|M)P(M)}{P(\mathcal{D})}$$

The *Bayes factor* allows the plausibility of two models (M_1 and M_2) to be compared:

$$K = \frac{p(\mathcal{D}|M_1)}{p(\mathcal{D}|M_2)} = \frac{\int_{\theta_{M_1}} p(\mathcal{D}|\theta_{M_1}, M_1)p(\theta_{M_1}|M_1)d\theta_{M_1}}{\int_{\theta_{M_2}} p(\mathcal{D}|\theta_{M_2}, M_2)p(\theta_{M_2}|M_2)d\theta_{M_2}}$$

This is usually too costly in practice, so approximations are used.

MODEL SELECTION

Some approximations/alternatives to the Bayesian approach:

- **Laplace approximations:** find the MAP/ML solution and use a Gaussian approximation to the parameter uncertainty.
- **Minimum Message Length (MML):** an information theoretic approach.
- **Minimum Description Length (MDL):** an information theoretic approach based on how well the model compresses the data.
- **Akaike Information Criterion (AIC):** $-2 \log p(\mathcal{D}|\theta) + 2k$, where k is the number of estimated parameters.
- **Bayesian Information Criterion (BIC):**
 $-2 \log p(\mathcal{D}|\theta) + k \log q$, where q is the number of observations.

ACCURACY MEASURES FOR REGRESSION

- **Root-mean squared error** for point predictions.
- **Correlation coefficient** for point predictions.
- **Log predictive probability** can be used for probabilistic predictions.
- **Expected loss/risk** for point predictions for decision making.

ACCURACY MEASURES FOR BINARY CLASSIFICATION


		Condition (as determined by "Gold standard")			
		Total population	Condition positive	Condition negative	Prevalence = $\frac{\Sigma \text{ Condition positive}}{\Sigma \text{ Total population}}$
Test outcome	Test outcome positive	True positive	False positive (Type I error)	Positive predictive value (PPV, Precision) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Test outcome positive}}$	False discovery rate (FDR) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Test outcome positive}}$
	Test outcome negative	False negative (Type II error)	True negative	False omission rate (FOR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Test outcome negative}}$	Negative predictive value (NPV) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Test outcome negative}}$
		Positive likelihood ratio (LR+) = $\frac{\text{TPR/FPR}}{\text{TPR/FPR}}$	True positive rate (TPR, Sensitivity, Recall) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Condition positive}}$	False positive rate (FPR, Fall-out) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Condition negative}}$	Accuracy (ACC) = $\frac{\Sigma \text{ True positive} + \Sigma \text{ True negative}}{\Sigma \text{ Total population}}$
		Negative likelihood ratio (LR-) = $\frac{\text{FNR/TNR}}{\text{FNR/TNR}}$	False negative rate (FNR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Condition positive}}$	True negative rate (TNR, Specificity, SPC) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Condition negative}}$	
		Diagnostic odds ratio (DOR) = $\frac{\text{LR+}}{\text{LR-}}$			

Wikipedia contributors,
"Sensitivity and specificity,"
Wikipedia, The Free
Encyclopedia, http://en.wikipedia.org/w/index.php?title=Sensitivity_and_specificity&oldid=655245669
(accessed April 9, 2015).

ACCURACY MEASURES FROM ROC CURVE

The **Receiver operating characteristic** (ROC) curve is a plot of *true-positive rate* (sensitivity) versus *false-positive rate* (1-specificity) over the full range of possible thresholds.

The **area under the curve** (AUC) is the integral under the ROC curve.



LOG PREDICTIVE PROBABILITY

Some data are more easily classified than others.
Probabilistic classifiers provide a level of confidence for each prediction.

$$p(y_* | \mathbf{x}_*, \mathbf{y}, \mathbf{X}, \theta)$$

Quality of predictions can be assessed using the **test log predictive probability**:

$$\frac{1}{m} \sum_{i=1}^m \log_2 p(y_{*i} = t_i | \mathbf{x}_{*i}, \mathbf{y}, \mathbf{X}, \theta)$$


After subtracting the baseline measure, this shows the average bits of information given by the model.

Rasmussen & Williams. "Gaussian Processes for Machine Learning", MIT Press (2006).

<http://www.gaussianprocess.org/gpml/>

GENERATIVE MODELS FOR CLASSIFICATION

$$P(y=k|\mathbf{x}) = \frac{P(y=k)p(\mathbf{x}|y=k)}{\sum_j P(y=j)p(\mathbf{x}|y=j)}$$



LINEAR DISCRIMINANT ANALYSIS





$$P(y=k|\mathbf{x}) = \frac{P(y=k)p(\mathbf{x}|y=k)}{\sum_j P(y=j)p(\mathbf{x}|y=j)}$$

Assumes:

$$P(\mathbf{x}|y=k) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma})$$

Model has $2p + p(p-1)$ parameters to estimate (two means and a single covariance).

Number of observations is pn (size of inputs).







QUADRATIC DISCRIMINANT ANALYSIS

$$P(y=k|\mathbf{x}) = \frac{P(y=k)p(\mathbf{x}|y=k)}{\sum_j P(y=j)p(\mathbf{x}|y=j)}$$

Assumes different covariances:

$$P(\mathbf{x}|y=k) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$



Model has $2p + 2p(p-1)$ parameters to estimate (two means and two covariances).

Number of observations is pn .

NAIVE BAYES





$$P(y=k|\mathbf{x}) = \frac{P(y=k)p(\mathbf{x}|y=k)}{\sum_j P(y=j)p(\mathbf{x}|y=j)}$$

Assumes that features are independent:

$$p(\mathbf{x}|y=k) = \prod_i p(x_i|y=k)$$

Model has variable number of parameters to estimate, but the above example has $3p$.

Number of observations is pn .



LINEAR REGRESSION: MAXIMUM LIKELIHOOD

A simple way to do regression is by:

$$f(\mathbf{x}_*) = \mathbf{a}^T \mathbf{x}_*$$

Assuming Gaussian noise on \mathbf{y} , the ML estimate of \mathbf{a} is by:

$$\hat{\mathbf{a}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where

$$\mathbf{X} = (\mathbf{x}_1 \quad \mathbf{x}_2 \quad \dots \quad \mathbf{x}_n)^T, \text{ and } \mathbf{y} = (y_1 \quad y_2 \quad \dots y_n)^T$$

Model has p parameters to estimate.

Number of observations is n (number of targets).

Usually needs dimensionality reduction, with (eg) SVD.

LINEAR REGRESSION: MAXIMUM POSTERIOR

We may have prior knowledge about various distributions:

$$p(y_*|\mathbf{x}_*, \mathbf{a}) = \mathcal{N}(\mathbf{a}^T \mathbf{x}_*, \sigma^2)$$
$$p(\mathbf{a}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma}_0)$$

Therefore,

$$p(\mathbf{a}|\mathbf{y}, \mathbf{X}) = \mathcal{N}(\sigma^{-2} \mathbf{B}^{-1} \mathbf{X}^T \mathbf{y}, \mathbf{B}^{-1}), \text{ where } \mathbf{B} = \sigma^{-2} \mathbf{X}^T \mathbf{X} + \mathbf{\Sigma}_0^{-1}$$

Maximum a posteriori (MAP) estimate of \mathbf{a} is by:

$$\hat{\mathbf{a}} = \sigma^{-2} \mathbf{B}^{-1} \mathbf{X}^T \mathbf{y}, \text{ where } \mathbf{B} = \sigma^{-2} \mathbf{X}^T \mathbf{X} + \mathbf{\Sigma}_0^{-1}$$

LINEAR REGRESSION: BAYESIAN

We may have prior knowledge about various distributions:

$$p(y_*|\mathbf{x}_*, \mathbf{a}) = \mathcal{N}(\mathbf{a}^T \mathbf{x}_*, \sigma^2)$$

$$p(\mathbf{a}) = \mathcal{N}(\mathbf{0}, \mathbf{\Sigma}_0)$$

Therefore,

$$p(\mathbf{a}|\mathbf{y}, \mathbf{X}) = \mathcal{N}(\sigma^{-2} \mathbf{B}^{-1} \mathbf{X}^T \mathbf{y}, \mathbf{B}^{-1}), \text{ where } \mathbf{B} = \sigma^{-2} \mathbf{X}^T \mathbf{X} + \mathbf{\Sigma}_0^{-1}$$

Predictions are made by integrating out the uncertainty of the weights, rather than estimating them:

$$p(y_*|\mathbf{x}_*, \mathbf{y}, \mathbf{X}) = \int_{\mathbf{a}} p(y_*|\mathbf{x}_*, \mathbf{a}) p(\mathbf{a}|\mathbf{y}, \mathbf{X}) d\mathbf{a}$$

$$= \mathcal{N}(\sigma^{-2} \mathbf{x}_*^T \mathbf{B}^{-1} \mathbf{X}^T \mathbf{y}, \mathbf{x}_*^T \mathbf{B}^{-1} \mathbf{x}_*)$$

Estimated parameters may be σ^2 , and parameters encoding $\mathbf{\Sigma}_0$.

KERNEL METHODS: WOODBURY MATRIX IDENTITY

$$\begin{aligned}\mathbf{B}^{-1} &= \left(\sigma^{-2} \mathbf{X}^T \mathbf{X} + \boldsymbol{\Sigma}_0^{-1} \right)^{-1} \\ &= \boldsymbol{\Sigma}_0 - \boldsymbol{\Sigma}_0 \mathbf{X}^T (\mathbf{I} \sigma^2 + \mathbf{X} \boldsymbol{\Sigma}_0 \mathbf{X}^T)^{-1} \mathbf{X} \boldsymbol{\Sigma}_0\end{aligned}$$

Wikipedia contributors, "Woodbury matrix identity," Wikipedia, The Free Encyclopedia, http://en.wikipedia.org/w/index.php?title=Woodbury_matrix_identity&oldid=638370219 (accessed April 1, 2015).

Dimensions of $\mathbf{X}^T \mathbf{X}$ are $p \times p$.

Dimensions of $\mathbf{X} \boldsymbol{\Sigma}_0 \mathbf{X}^T$ are $n \times n$.

KERNEL METHODS: GAUSSIAN PROCESS REGRESSION

The predicted distribution is:

$$p(y_* | \mathbf{x}_*, \mathbf{y}, \mathbf{X}) = \mathcal{N}(\mathbf{k}^T \mathbf{C}^{-1} \mathbf{y}, c - \mathbf{k}^T \mathbf{C}^{-1} \mathbf{k})$$

where:

$$\mathbf{C} = \mathbf{X} \mathbf{\Sigma}_0 \mathbf{X}^T + \mathbf{I} \sigma^2$$

$$\mathbf{k} = \mathbf{X} \mathbf{\Sigma}_0 \mathbf{x}_*$$

$$c = \mathbf{x}_*^T \mathbf{\Sigma}_0 \mathbf{x}_* + \sigma^2$$

KERNEL METHODS: NONLINEAR METHODS

Sometimes, we want alternatives to $\mathbf{C} = \mathbf{X}\mathbf{\Sigma}_0\mathbf{X}^T + \mathbf{I}\sigma^2$.

Nonlinearity is achieved by replacing the matrix $\mathbf{K} = \mathbf{X}\mathbf{\Sigma}_0\mathbf{X}^T$ with some function of the data that gives a positive definite matrix encoding similarities.

eg

$$k(\mathbf{x}_i, \mathbf{x}_j) = \theta_1 + \theta_2 \mathbf{x}_i \cdot \mathbf{x}_j + \theta_3 \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\theta_4}\right)$$

Hyper-parameters θ_1 to θ_4 can be optimised in a number of ways.

KERNEL METHODS: NONLINEAR METHODS

For large p , small n problems, nonlinear methods do not seem to help much.


Nonlinearity also reduces interpretability.

DISCRIMINATIVE MODELS FOR CLASSIFICATION

$$t = \sigma(f(\mathbf{x}_*))$$

where σ is some squashing function, eg:

- Heaviside step function.
- Logistic function (inverse of Logit).
- Normal CDF (inverse of Probit).




PROBABILISTIC CLASSIFICATION


Integrating over the uncertainty of the separating hyperplane allows probabilistic predictions further from the training data. This is not usually done for methods such as the relevance-vector machine (RVM).

Rasmussen, Carl Edward, and Joaquin Quinonero-Candela. "Healing the relevance vector machine through augmentation." In Proceedings of the 22nd international conference on Machine learning, pp. 689-696. ACM, 2005.


Simple Logistic Regression




Hyperplane Uncertainty



Bayesian Logistic Regression



Bayesian Logistic Regression



PROBABILISTIC CLASSIFICATION

Making probabilistic predictions involves:

- 1 Computing the distribution of a latent variable corresponding to the test data (cf regression):

$$p(f_*|\mathbf{x}_*, \mathbf{y}, \mathbf{X}) = \int_{\mathbf{f}} p(f_*|\mathbf{x}_*, \mathbf{f})p(\mathbf{f}|\mathbf{y}, \mathbf{X})d\mathbf{f}$$

- 2 Using this distribution to give a probabilistic prediction:


$$P(y_* = 1|\mathbf{x}_*, \mathbf{y}, \mathbf{X}) = \int_{f_*} \sigma(f_*)p(f_*|\mathbf{x}_*, \mathbf{y}, \mathbf{X})df_*$$

Unfortunately, the second integral is analytically intractable, so approximations are needed.

PROBABILISTIC CLASSIFICATION

Approximate methods for probabilistic classification include:


- **The Laplace Approximation (LA)**. Fastest, but less accurate.
- **Expectation Propagation (EP)**. More accurate than the Laplace approximation, but slightly slower.
- **MCMC** methods. The “gold standard”, but very slow to draw lots of random samples.



Nickisch, Hannes, and Carl Edward Rasmussen. “Approximations for Binary Gaussian Process Classification.”
Journal of Machine Learning Research 9 (2008): 2035-2078.

SUPPORT VECTOR CLASSIFICATION

If you are only interested in binary predictions, support-vector machines are reasonably fast and accurate.



ENSEMBLE LEARNING

Combining predictions from weak learners.

- **Bootstrap aggregating (bagging)**

- Train several weak classifiers, with different models or randomly drawn subsets of the data.
- Average their predictions with equal weight.

- **Boosting**

- A family of approaches, where models are weighted according to their accuracy.
- AdaBoost is popular, but has problems with target noise.

- **Bayesian model averaging**

- Really a model selection method.
- Relatively ineffective for combining models.

- **Bayesian model combination**

- Shows promise.

Monteith, et al. "Turning Bayesian model averaging into Bayesian model combination." Neural Networks (IJCNN), The 2011 International Joint Conference on. IEEE, 2011.